Table of Contents

1 Introduction **1**

1.1 Language Processors 1
1.1.1 Exercises for Section 1.1 3
1.2 The Structure of a Compiler 4
1.2.1 Lexical Analysis 5
1.2.2 Syntax Analysis 8
1.2.3 Semantic Analysis 8
1.2.4 Intermediate Code Generation 9
1.2.5 Code Optimization 10
1.2.6 Code Generation 10
1.2.7 Symbol-Table Management 11
1.2.8 The Grouping of Phases into Passes 11
1.2.9 Compiler-Construction Tools 12
1.3 The Evolution of Programming Languages 12
1.3.1 The Move to Higher-level Languages 13
1.3.2 Impacts on Compilers 14
1.3.3 Exercises for Section 1.3 14
1.4 The Science of Building a Compiler 15
1.4.1 Modeling in Compiler Design and Implementation 15
1.4.2 The Science of Code Optimization 15
1.5 Applications of Compiler Technology 17
1.5.1 Implementation of High-Level Programming Languages . 17
1.5.2 Optimizations for Computer Architectures 19
1.5.3 Design of New Computer Architectures 21
1.5.4 Program Translations 22
1.5.5 Software Productivity Tools 23
1.6 Programming Language Basics 25
1.6.1 The Static/Dynamic Distinction 25
1.6.2 Environments and States 26
1.6.3 Static Scope and Block Structure 28
1.6.4 Explicit Access Control 31
1.6.5 Dynamic Scope 31
1.6.6 Parameter Passing Mechanisms 33

1.6.7 Aliasing 35
1.6.8 Exercises for Section 1.6 35
1.7 Summary of Chapter 1 36
1.8 References for Chapter 1 38

2 A Simple Syntax-Directed Translator **39**
2.1 Introduction 40
2.2 Syntax Definition 42
2.2.1 Definition of Grammars 42
2.2.2 Derivations 44
2.2.3 Parse Trees 45
2.2.4 Ambiguity 47
2.2.5 Associativity of Operators 48
2.2.6 Precedence of Operators 48
2.2.7 Exercises for Section 2.2 51
2.3 Syntax-Directed Translation 52
2.3.1 Postfix Notation 53
2.3.2 Synthesized Attributes 54
2.3.3 Simple Syntax-Directed Definitions 56
2.3.4 Tree Traversals 56
2.3.5 Translation Schemes 57
2.3.6 Exercises for Section 2.3 60
2.4 Parsing 60
2.4.1 Top-Down Parsing 61
2.4.2 Predictive Parsing 64
2.4.3 When to Use ϵ-Productions 65
2.4.4 Designing a Predictive Parser 66
2.4.5 Left Recursion 67
2.4.6 Exercises for Section 2.4 68
2.5 A Translator for Simple Expressions 68
2.5.1 Abstract and Concrete Syntax 69
2.5.2 Adapting the Translation Scheme 70
2.5.3 Procedures for the Nonterminals 72
2.5.4 Simplifying the Translator 73
2.5.5 The Complete Program 74
2.6 Lexical Analysis 76
2.6.1 Removal of White Space and Comments 77
2.6.2 Reading Ahead 78
2.6.3 Constants 78
2.6.4 Recognizing Keywords and Identifiers 79
2.6.5 A Lexical Analyzer 81
2.6.6 Exercises for Section 2.6 84
2.7 Symbol Tables 85
2.7.1 Symbol Table Per Scope 86
2.7.2 The Use of Symbol Tables 89

2.8 Intermediate Code Generation 91
2.8.1 Two Kinds of Intermediate Representations 91
2.8.2 Construction of Syntax Trees 92
2.8.3 Static Checking 97
2.8.4 Three-Address Code 99
2.8.5 Exercises for Section 2.8 105
2.9 Summary of Chapter 2 105

3 Lexical Analysis **109**
3.1 The Role of the Lexical Analyzer 109
3.1.1 Lexical Analysis Versus Parsing 110
3.1.2 Tokens, Patterns, and Lexemes 111
3.1.3 Attributes for Tokens 112
3.1.4 Lexical Errors 113
3.1.5 Exercises for Section 3.1 114
3.2 Input Buffering 115
3.2.1 Buffer Pairs 115
3.2.2 Sentinels 116
3.3 Specification of Tokens 116
3.3.1 Strings and Languages 117
3.3.2 Operations on Languages 119
3.3.3 Regular Expressions 120
3.3.4 Regular Definitions 123
3.3.5 Extensions of Regular Expressions 124
3.3.6 Exercises for Section 3.3 125
3.4 Recognition of Tokens 128
3.4.1 Transition Diagrams 130
3.4.2 Recognition of Reserved Words and Identifiers 132
3.4.3 Completion of the Running Example 133
3.4.4 Architecture of a Transition-Diagram-Based Lexical Analyzer 134
3.4.5 Exercises for Section 3.4 136
3.5 The Lexical-Analyzer Generator `Lex` 140
3.5.1 Use of `Lex` 140
3.5.2 Structure of `Lex` Programs 141
3.5.3 Conflict Resolution in `Lex` 144
3.5.4 The Lookahead Operator 144
3.5.5 Exercises for Section 3.5 146
3.6 Finite Automata 147
3.6.1 Nondeterministic Finite Automata 147
3.6.2 Transition Tables 148
3.6.3 Acceptance of Input Strings by Automata 149
3.6.4 Deterministic Finite Automata 149
3.6.5 Exercises for Section 3.6 151
3.7 From Regular Expressions to Automata 152

3.7.1 Conversion of an NFA to a DFA . . . 152
3.7.2 Simulation of an NFA . . . 156
3.7.3 Efficiency of NFA Simulation . . . 157
3.7.4 Construction of an NFA from a Regular Expression . . . 159
3.7.5 Efficiency of String-Processing Algorithms . . . 163
3.7.6 Exercises for Section 3.7 . . . 166
3.8 Design of a Lexical-Analyzer Generator . . . 166
3.8.1 The Structure of the Generated Analyzer . . . 167
3.8.2 Pattern Matching Based on NFA's . . . 168
3.8.3 DFA's for Lexical Analyzers . . . 170
3.8.4 Implementing the Lookahead Operator . . . 171
3.8.5 Exercises for Section 3.8 . . . 172
3.9 Optimization of DFA-Based Pattern Matchers . . . 173
3.9.1 Important States of an NFA . . . 173
3.9.2 Functions Computed From the Syntax Tree . . . 175
3.9.3 Computing *nullable*, *firstpos*, and *lastpos* . . . 176
3.9.4 Computing *followpos* . . . 177
3.9.5 Converting a Regular Expression Directly to a DFA . . . 179
3.9.6 Minimizing the Number of States of a DFA . . . 180
3.9.7 State Minimization in Lexical Analyzers . . . 184
3.9.8 Trading Time for Space in DFA Simulation . . . 185
3.9.9 Exercises for Section 3.9 . . . 186
3.10 Summary of Chapter 3 . . . 187
3.11 References for Chapter 3 . . . 189

4 Syntax Analysis **191**
4.1 Introduction . . . 192
4.1.1 The Role of the Parser . . . 192
4.1.2 Representative Grammars . . . 193
4.1.3 Syntax Error Handling . . . 194
4.1.4 Error-Recovery Strategies . . . 195
4.2 Context-Free Grammars . . . 197
4.2.1 The Formal Definition of a Context-Free Grammar . . . 197
4.2.2 Notational Conventions . . . 198
4.2.3 Derivations . . . 199
4.2.4 Parse Trees and Derivations . . . 201
4.2.5 Ambiguity . . . 203
4.2.6 Verifying the Language Generated by a Grammar . . . 204
4.2.7 Context-Free Grammars Versus Regular Expressions . . . 205
4.2.8 Exercises for Section 4.2 . . . 206
4.3 Writing a Grammar . . . 209
4.3.1 Lexical Versus Syntactic Analysis . . . 209
4.3.2 Eliminating Ambiguity . . . 210
4.3.3 Elimination of Left Recursion . . . 212
4.3.4 Left Factoring . . . 214

4.3.5 Non-Context-Free Language Constructs . . . 215
4.3.6 Exercises for Section 4.3 . . . 216
4.4 Top-Down Parsing . . . 217
4.4.1 Recursive-Descent Parsing . . . 219
4.4.2 FIRST and FOLLOW . . . 220
4.4.3 LL(1) Grammars . . . 222
4.4.4 Nonrecursive Predictive Parsing . . . 226
4.4.5 Error Recovery in Predictive Parsing . . . 228
4.4.6 Exercises for Section 4.4 . . . 231
4.5 Bottom-Up Parsing . . . 233
4.5.1 Reductions . . . 234
4.5.2 Handle Pruning . . . 235
4.5.3 Shift-Reduce Parsing . . . 236
4.5.4 Conflicts During Shift-Reduce Parsing . . . 238
4.5.5 Exercises for Section 4.5 . . . 240
4.6 Introduction to LR Parsing: Simple LR . . . 241
4.6.1 Why LR Parsers? . . . 241
4.6.2 Items and the LR(0) Automaton . . . 242
4.6.3 The LR-Parsing Algorithm . . . 248
4.6.4 Constructing SLR-Parsing Tables . . . 252
4.6.5 Viable Prefixes . . . 256
4.6.6 Exercises for Section 4.6 . . . 257
4.7 More Powerful LR Parsers . . . 259
4.7.1 Canonical LR(1) Items . . . 260
4.7.2 Constructing LR(1) Sets of Items . . . 261
4.7.3 Canonical LR(1) Parsing Tables . . . 265
4.7.4 Constructing LALR Parsing Tables . . . 266
4.7.5 Efficient Construction of LALR Parsing Tables . . . 270
4.7.6 Compaction of LR Parsing Tables . . . 275
4.7.7 Exercises for Section 4.7 . . . 277
4.8 Using Ambiguous Grammars . . . 278
4.8.1 Precedence and Associativity to Resolve Conflicts . . . 279
4.8.2 The "Dangling-Else" Ambiguity . . . 281
4.8.3 Error Recovery in LR Parsing . . . 283
4.8.4 Exercises for Section 4.8 . . . 285
4.9 Parser Generators . . . 287
4.9.1 The Parser Generator `Yacc` . . . 287
4.9.2 Using `Yacc` with Ambiguous Grammars . . . 291
4.9.3 Creating `Yacc` Lexical Analyzers with `Lex` . . . 294
4.9.4 Error Recovery in `Yacc` . . . 295
4.9.5 Exercises for Section 4.9 . . . 297
4.10 Summary of Chapter 4 . . . 297
4.11 References for Chapter 4 . . . 300

5 Syntax-Directed Translation **303**
5.1 Syntax-Directed Definitions 304
5.1.1 Inherited and Synthesized Attributes 304
5.1.2 Evaluating an SDD at the Nodes of a Parse Tree 306
5.1.3 Exercises for Section 5.1 309
5.2 Evaluation Orders for SDD's 310
5.2.1 Dependency Graphs 310
5.2.2 Ordering the Evaluation of Attributes 312
5.2.3 S-Attributed Definitions 312
5.2.4 L-Attributed Definitions 313
5.2.5 Semantic Rules with Controlled Side Effects 314
5.2.6 Exercises for Section 5.2 317
5.3 Applications of Syntax-Directed Translation 318
5.3.1 Construction of Syntax Trees 318
5.3.2 The Structure of a Type 321
5.3.3 Exercises for Section 5.3 323
5.4 Syntax-Directed Translation Schemes 324
5.4.1 Postfix Translation Schemes 324
5.4.2 Parser-Stack Implementation of Postfix SDT's 325
5.4.3 SDT's With Actions Inside Productions 327
5.4.4 Eliminating Left Recursion From SDT's 328
5.4.5 SDT's for L-Attributed Definitions 331
5.4.6 Exercises for Section 5.4 336
5.5 Implementing L-Attributed SDD's 337
5.5.1 Translation During Recursive-Descent Parsing 338
5.5.2 On-The-Fly Code Generation 340
5.5.3 L-Attributed SDD's and LL Parsing 343
5.5.4 Bottom-Up Parsing of L-Attributed SDD's 348
5.5.5 Exercises for Section 5.5 352
5.6 Summary of Chapter 5 353
5.7 References for Chapter 5 354

6 Intermediate-Code Generation **357**
6.1 Variants of Syntax Trees 358
6.1.1 Directed Acyclic Graphs for Expressions 359
6.1.2 The Value-Number Method for Constructing DAG's . . . 360
6.1.3 Exercises for Section 6.1 362
6.2 Three-Address Code 363
6.2.1 Addresses and Instructions 364
6.2.2 Quadruples 366
6.2.3 Triples 367
6.2.4 Static Single-Assignment Form 369
6.2.5 Exercises for Section 6.2 370
6.3 Types and Declarations 370
6.3.1 Type Expressions 371

6.3.2 Type Equivalence . . . 372
6.3.3 Declarations . . . 373
6.3.4 Storage Layout for Local Names . . . 373
6.3.5 Sequences of Declarations . . . 376
6.3.6 Fields in Records and Classes . . . 376
6.3.7 Exercises for Section 6.3 . . . 378
6.4 Translation of Expressions . . . 378
6.4.1 Operations Within Expressions . . . 378
6.4.2 Incremental Translation . . . 380
6.4.3 Addressing Array Elements . . . 381
6.4.4 Translation of Array References . . . 383
6.4.5 Exercises for Section 6.4 . . . 384
6.5 Type Checking . . . 386
6.5.1 Rules for Type Checking . . . 387
6.5.2 Type Conversions . . . 388
6.5.3 Overloading of Functions and Operators . . . 390
6.5.4 Type Inference and Polymorphic Functions . . . 391
6.5.5 An Algorithm for Unification . . . 395
6.5.6 Exercises for Section 6.5 . . . 398
6.6 Control Flow . . . 399
6.6.1 Boolean Expressions . . . 399
6.6.2 Short-Circuit Code . . . 400
6.6.3 Flow-of-Control Statements . . . 401
6.6.4 Control-Flow Translation of Boolean Expressions . . . 403
6.6.5 Avoiding Redundant Gotos . . . 405
6.6.6 Boolean Values and Jumping Code . . . 408
6.6.7 Exercises for Section 6.6 . . . 408
6.7 Backpatching . . . 410
6.7.1 One-Pass Code Generation Using Backpatching . . . 410
6.7.2 Backpatching for Boolean Expressions . . . 411
6.7.3 Flow-of-Control Statements . . . 413
6.7.4 Break-, Continue-, and Goto-Statements . . . 416
6.7.5 Exercises for Section 6.7 . . . 417
6.8 Switch-Statements . . . 418
6.8.1 Translation of Switch-Statements . . . 419
6.8.2 Syntax-Directed Translation of Switch-Statements . . . 420
6.8.3 Exercises for Section 6.8 . . . 421
6.9 Intermediate Code for Procedures . . . 422
6.10 Summary of Chapter 6 . . . 424
6.11 References for Chapter 6 . . . 425

7 Run-Time Environments 427
7.1 Storage Organization . . . 427
7.1.1 Static Versus Dynamic Storage Allocation . . . 429
7.2 Stack Allocation of Space . . . 430
7.2.1 Activation Trees . . . 430
7.2.2 Activation Records . . . 433
7.2.3 Calling Sequences . . . 436
7.2.4 Variable-Length Data on the Stack . . . 438
7.2.5 Exercises for Section 7.2 . . . 440
7.3 Access to Nonlocal Data on the Stack . . . 441
7.3.1 Data Access Without Nested Procedures . . . 442
7.3.2 Issues With Nested Procedures . . . 442
7.3.3 A Language With Nested Procedure Declarations . . . 443
7.3.4 Nesting Depth . . . 443
7.3.5 Access Links . . . 445
7.3.6 Manipulating Access Links . . . 447
7.3.7 Access Links for Procedure Parameters . . . 448
7.3.8 Displays . . . 449
7.3.9 Exercises for Section 7.3 . . . 451
7.4 Heap Management . . . 452
7.4.1 The Memory Manager . . . 453
7.4.2 The Memory Hierarchy of a Computer . . . 454
7.4.3 Locality in Programs . . . 455
7.4.4 Reducing Fragmentation . . . 457
7.4.5 Manual Deallocation Requests . . . 460
7.4.6 Exercises for Section 7.4 . . . 463
7.5 Introduction to Garbage Collection . . . 463
7.5.1 Design Goals for Garbage Collectors . . . 464
7.5.2 Reachability . . . 466
7.5.3 Reference Counting Garbage Collectors . . . 468
7.5.4 Exercises for Section 7.5 . . . 470
7.6 Introduction to Trace-Based Collection . . . 470
7.6.1 A Basic Mark-and-Sweep Collector . . . 471
7.6.2 Basic Abstraction . . . 473
7.6.3 Optimizing Mark-and-Sweep . . . 475
7.6.4 Mark-and-Compact Garbage Collectors . . . 476
7.6.5 Copying collectors . . . 478
7.6.6 Comparing Costs . . . 482
7.6.7 Exercises for Section 7.6 . . . 482
7.7 Short-Pause Garbage Collection . . . 483
7.7.1 Incremental Garbage Collection . . . 483
7.7.2 Incremental Reachability Analysis . . . 485
7.7.3 Partial-Collection Basics . . . 487
7.7.4 Generational Garbage Collection . . . 488
7.7.5 The Train Algorithm . . . 490

7.7.6 Exercises for Section 7.7 . . . 493
7.8 Advanced Topics in Garbage Collection . . . 494
7.8.1 Parallel and Concurrent Garbage Collection . . . 495
7.8.2 Partial Object Relocation . . . 497
7.8.3 Conservative Collection for Unsafe Languages . . . 498
7.8.4 Weak References . . . 498
7.8.5 Exercises for Section 7.8 . . . 499
7.9 Summary of Chapter 7 . . . 500
7.10 References for Chapter 7 . . . 502

8 Code Generation **505**
8.1 Issues in the Design of a Code Generator . . . 506
8.1.1 Input to the Code Generator . . . 507
8.1.2 The Target Program . . . 507
8.1.3 Instruction Selection . . . 508
8.1.4 Register Allocation . . . 510
8.1.5 Evaluation Order . . . 511
8.2 The Target Language . . . 512
8.2.1 A Simple Target Machine Model . . . 512
8.2.2 Program and Instruction Costs . . . 515
8.2.3 Exercises for Section 8.2 . . . 516
8.3 Addresses in the Target Code . . . 518
8.3.1 Static Allocation . . . 518
8.3.2 Stack Allocation . . . 520
8.3.3 Run-Time Addresses for Names . . . 522
8.3.4 Exercises for Section 8.3 . . . 524
8.4 Basic Blocks and Flow Graphs . . . 525
8.4.1 Basic Blocks . . . 526
8.4.2 Next-Use Information . . . 528
8.4.3 Flow Graphs . . . 529
8.4.4 Representation of Flow Graphs . . . 530
8.4.5 Loops . . . 531
8.4.6 Exercises for Section 8.4 . . . 531
8.5 Optimization of Basic Blocks . . . 533
8.5.1 The DAG Representation of Basic Blocks . . . 533
8.5.2 Finding Local Common Subexpressions . . . 534
8.5.3 Dead Code Elimination . . . 535
8.5.4 The Use of Algebraic Identities . . . 536
8.5.5 Representation of Array References . . . 537
8.5.6 Pointer Assignments and Procedure Calls . . . 539
8.5.7 Reassembling Basic Blocks From DAG's . . . 539
8.5.8 Exercises for Section 8.5 . . . 541
8.6 A Simple Code Generator . . . 542
8.6.1 Register and Address Descriptors . . . 543
8.6.2 The Code-Generation Algorithm . . . 544

8.6.3 Design of the Function *getReg* 547
8.6. 548
8.7 Peephole Optimization 549
8.7.1 Eliminating Redundant Loads and Stores 550
8.7.2 Eliminating Unreachable Code 550
8.7.3 Flow-of-Control Optimizations 551
8.7.4 Algebraic Simplification and Reduction in Strength 552
8.7.5 Use of Machine Idioms 552
8.7. Exercises for Section 8.7 553
8.8 Register Allocation and Assignment 553
8.8.1 Global Register Allocation 553
8.8.2 Usage Counts 554
8.8.3 Register Assignment for Outer Loops 556
8.8.4 Register Allocation by Graph Coloring 556
8.8.5 Exercises for Section 8.8 557
8.9 Instruction Selection by Tree Rewriting 558
8.9.1 Tree-Translation Schemes 558
8.9.2 Code Generation by Tiling an Input Tree 560
8.9.3 Pattern Matching by Parsing 563
8.9.4 Routines for Semantic Checking 565
8.9.5 General Tree Matching 565
8.9.6 Exercises for Section 8.9 567
8.10 Optimal Code Generation for Expressions 567
8.10.1 Ershov Numbers 567
8.10.2 Generating Code From Labeled Expression Trees 568
8.10.3 Evaluating Expressions with an Insufficient Supply of Registers 570
8.10.4 Exercises for Section 8.10 572
8.11 Dynamic Programming Code-Generation 573
8.11.1 Contiguous Evaluation 574
8.11.2 The Dynamic Programming Algorithm 575
8.11.3 Exercises for Section 8.11 577
8.12 Summary of Chapter 8 578
8.13 References for Chapter 8 579

9 Machine-Independent Optimizations **583**
9.1 The Principal Sources of Optimization 584
9.1.1 Causes of Redundancy 584
9.1.2 A Running Example: Quicksort 585
9.1.3 Semantics-Preserving Transformations 586
9.1.4 Global Common Subexpressions 588
9.1.5 Copy Propagation 590
9.1.6 Dead-Code Elimination 591
9.1.7 Code Motion 592
9.1.8 Induction Variables and Reduction in Strength 592

9.1.9 Exercises for Section 9.1 . . . 596
9.2 Introduction to Data-Flow Analysis . . . 597
9.2.1 The Data-Flow Abstraction . . . 597
9.2.2 The Data-Flow Analysis Schema . . . 599
9.2.3 Data-Flow Schemas on Basic Blocks . . . 600
9.2.4 Reaching Definitions . . . 601
9.2.5 Live-Variable Analysis . . . 608
9.2.6 Available Expressions . . . 610
9.2.7 Summary . . . 614
9.2.8 Exercises for Section 9.2 . . . 615
9.3 Foundations of Data-Flow Analysis . . . 618
9.3.1 Semilattices . . . 618
9.3.2 Transfer Functions . . . 623
9.3.3 The Iterative Algorithm for General Frameworks . . . 626
9.3.4 Meaning of a Data-Flow Solution . . . 628
9.3.5 Exercises for Section 9.3 . . . 631
9.4 Constant Propagation . . . 632
9.4.1 Data-Flow Values for the Constant-Propagation Framework . . . 633
9.4.2 The Meet for the Constant-Propagation Framework . . . 633
9.4.3 Transfer Functions for the Constant-Propagation Framework . . . 634
9.4.4 Monotonicity of the Constant-Propagation Framework . . . 635
9.4.5 Nondistributivity of the Constant-Propagation Framework 635
9.4.6 Interpretation of the Results . . . 637
9.4.7 Exercises for Section 9.4 . . . 637
9.5 Partial-Redundancy Elimination . . . 639
9.5.1 The Sources of Redundancy . . . 639
9.5.2 Can All Redundancy Be Eliminated? . . . 642
9.5.3 The Lazy-Code-Motion Problem . . . 644
9.5.4 Anticipation of Expressions . . . 645
9.5.5 The Lazy-Code-Motion Algorithm . . . 646
9.5.6 Exercises for Section 9.5 . . . 655
9.6 Loops in Flow Graphs . . . 655
9.6.1 Dominators . . . 656
9.6.2 Depth-First Ordering . . . 660
9.6.3 Edges in a Depth-First Spanning Tree . . . 661
9.6.4 Back Edges and Reducibility . . . 662
9.6.5 Depth of a Flow Graph . . . 665
9.6.6 Natural Loops . . . 665
9.6.7 Speed of Convergence of Iterative Data-Flow Algorithms . 667
9.6.8 Exercises for Section 9.6 . . . 669
9.7 Region-Based Analysis . . . 672
9.7.1 Regions . . . 672
9.7.2 Region Hierarchies for Reducible Flow Graphs . . . 673

9.7.3 Overview of a Region-Based Analysis 676
9.7.4 Necessary Assumptions About Transfer Functions 678
9.7.5 An Algorithm for Region-Based Analysis 680
9.7.6 Handling Nonreducible Flow Graphs 684
9.7.7 Exercises for Section 9.7 686
9.8 Symbolic Analysis 686
9.8.1 Affine Expressions of Reference Variables 687
9.8.2 Data-Flow Problem Formulation 689
9.8.3 Region-Based Symbolic Analysis 694
9.8.4 Exercises for Section 9.8 699
9.9 Summary of Chapter 9 700
9.10 References for Chapter 9 703

10 Instruction-Level Parallelism **707**
10.1 Processor Architectures 708
10.1.1 Instruction Pipelines and Branch Delays 708
10.1.2 Pipelined Execution 709
10.1.3 Multiple Instruction Issue 710
10.2 Code-Scheduling Constraints 710
10.2.1 Data Dependence 711
10.2.2 Finding Dependences Among Memory Accesses 712
10.2.3 Tradeoff Between Register Usage and Parallelism 713
10.2.4 Phase Ordering Between Register Allocation and Code Scheduling 716
10.2.5 Control Dependence 716
10.2.6 Speculative Execution Support 717
10.2.7 A Basic Machine Model 719
10.2.8 Exercises for Section 10.2 720
10.3 Basic-Block Scheduling 721
10.3.1 Data-Dependence Graphs 722
10.3.2 List Scheduling of Basic Blocks 723
10.3.3 Prioritized Topological Orders 725
10.3.4 Exercises for Section 10.3 726
10.4 Global Code Scheduling 727
10.4.1 Primitive Code Motion 728
10.4.2 Upward Code Motion 730
10.4.3 Downward Code Motion 731
10.4.4 Updating Data Dependences 732
10.4.5 Global Scheduling Algorithms 732
10.4.6 Advanced Code Motion Techniques 736
10.4.7 Interaction with Dynamic Schedulers 737
10.4.8 Exercises for Section 10.4 737
10.5 Software Pipelining 738
10.5.1 Introduction 738
10.5.2 Software Pipelining of Loops 740

10.5.3 Register Allocation and Code Generation 743
10.5.4 Do-Across Loops 743
10.5.5 Goals and Constraints of Software Pipelining 745
10.5.6 A Software-Pipelining Algorithm 749
10.5.7 Scheduling Acyclic Data-Dependence Graphs 749
10.5.8 Scheduling Cyclic Dependence Graphs 751
10.5.9 Improvements to the Pipelining Algorithms 758
10.5.10 Modular Variable Expansion 758
10.5.11 Conditional Statements 761
10.5.12 Hardware Support for Software Pipelining 762
10.5.13 Exercises for Section 10.5 763
10.6 Summary of Chapter 10 765
10.7 References for Chapter 10 766

11 Optimizing for Parallelism and Locality 769
11.1 Basic Concepts 771
11.1.1 Multiprocessors 772
11.1.2 Parallelism in Applications 773
11.1.3 Loop-Level Parallelism 775
11.1.4 Data Locality 777
11.1.5 Introduction to Affine Transform Theory 778
11.2 Matrix Multiply: An In-Depth Example 782
11.2.1 The Matrix-Multiplication Algorithm 782
11.2.2 Optimizations 785
11.2.3 Cache Interference 788
11.2.4 Exercises for Section 11.2 788
11.3 Iteration Spaces 788
11.3.1 Constructing Iteration Spaces from Loop Nests 788
11.3.2 Execution Order for Loop Nests 791
11.3.3 Matrix Formulation of Inequalities 791
11.3.4 Incorporating Symbolic Constants 793
11.3.5 Controlling the Order of Execution 793
11.3.6 Changing Axes 798
11.3.7 Exercises for Section 11.3 799
11.4 Affine Array Indexes 801
11.4.1 Affine Accesses 802
11.4.2 Affine and Nonaffine Accesses in Practice 803
11.4.3 Exercises for Section 11.4 804
11.5 Data Reuse 804
11.5.1 Types of Reuse 805
11.5.2 Self Reuse 806
11.5.3 Self-Spatial Reuse 809
11.5.4 Group Reuse 811
11.5.5 Exercises for Section 11.5 814
11.6 Array Data-Dependence Analysis 815

11.6.1 Definition of Data Dependence of Array Accesses 816
11.6.2 Integer Linear Programming 817
11.6.3 The GCD Test 818
11.6.4 Heuristics for Solving Integer Linear Programs 820
11.6.5 Solving General Integer Linear Programs 823
11.6.6 Summary 825
11.6.7 Exercises for Section 11.6 826
11.7 Finding Synchronization-Free Parallelism 828
11.7.1 An Introductory Example 828
11.7.2 Affine Space Partitions 830
11.7.3 Space-Partition Constraints 831
11.7.4 Solving Space-Partition Constraints 835
11.7.5 A Simple Code-Generation Algorithm 838
11.7.6 Eliminating Empty Iterations 841
11.7.7 Eliminating Tests from Innermost Loops 844
11.7.8 Source-Code Transforms 846
11.7.9 Exercises for Section 11.7 851
11.8 Synchronization Between Parallel Loops 853
11.8.1 A Constant Number of Synchronizations 853
11.8.2 Program-Dependence Graphs 854
11.8.3 Hierarchical Time 857
11.8.4 The Parallelization Algorithm 859
11.8.5 Exercises for Section 11.8 860
11.9 Pipelining 861
11.9.1 What is Pipelining? 861
11.9.2 Successive Over-Relaxation (SOR): An Example 863
11.9.3 Fully Permutable Loops 864
11.9.4 Pipelining Fully Permutable Loops 864
11.9.5 General Theory 867
11.9.6 Time-Partition Constraints 868
11.9.7 Solving Time-Partition Constraints by Farkas' Lemma . . 872
11.9.8 Code Transformations 875
11.9.9 Parallelism With Minimum Synchronization 880
11.9.10 Exercises for Section 11.9 882
11.10 Locality Optimizations 884
11.10.1 Temporal Locality of Computed Data 885
11.10.2 Array Contraction 885
11.10.3 Partition Interleaving 887
11.10.4 Putting it All Together 890
11.10.5 Exercises for Section 11.10 892
11.11 Other Uses of Affine Transforms 893
11.11.1 Distributed memory machines 894
11.11.2 Multi-Instruction-Issue Processors 895
11.11.3 Vector and SIMD Instructions 895
11.11.4 Prefetching 896

11.12 Summary of Chapter 11 . . . 897
11.13 References for Chapter 11 . . . 899

12 Interprocedural Analysis 903
12.1 Basic Concepts . . . 904
12.1.1 Call Graphs . . . 904
12.1.2 Context Sensitivity . . . 906
12.1.3 Call Strings . . . 908
12.1.4 Cloning-Based Context-Sensitive Analysis . . . 910
12.1.5 Summary-Based Context-Sensitive Analysis . . . 911
12.1.6 Exercises for Section 12.1 . . . 914
12.2 Why Interprocedural Analysis? . . . 916
12.2.1 Virtual Method Invocation . . . 916
12.2.2 Pointer Alias Analysis . . . 917
12.2.3 Parallelization . . . 917
12.2.4 Detection of Software Errors and Vulnerabilities . . . 917
12.2.5 SQL Injection . . . 918
12.2.6 Buffer Overflow . . . 920
12.3 A Logical Representation of Data Flow . . . 921
12.3.1 Introduction to Datalog . . . 921
12.3.2 Datalog Rules . . . 922
12.3.3 Intensional and Extensional Predicates . . . 924
12.3.4 Execution of Datalog Programs . . . 927
12.3.5 Incremental Evaluation of Datalog Programs . . . 928
12.3.6 Problematic Datalog Rules . . . 930
12.3.7 Exercises for Section 12.3 . . . 932
12.4 A Simple Pointer-Analysis Algorithm . . . 933
12.4.1 Why is Pointer Analysis Difficult . . . 934
12.4.2 A Model for Pointers and References . . . 935
12.4.3 Flow Insensitivity . . . 936
12.4.4 The Formulation in Datalog . . . 937
12.4.5 Using Type Information . . . 938
12.4.6 Exercises for Section 12.4 . . . 939
12.5 Context-Insensitive Interprocedural Analysis . . . 941
12.5.1 Effects of a Method Invocation . . . 941
12.5.2 Call Graph Discovery in Datalog . . . 943
12.5.3 Dynamic Loading and Reflection . . . 944
12.5.4 Exercises for Section 12.5 . . . 945
12.6 Context-Sensitive Pointer Analysis . . . 945
12.6.1 Contexts and Call Strings . . . 946
12.6.2 Adding Context to Datalog Rules . . . 949
12.6.3 Additional Observations About Sensitivity . . . 949
12.6.4 Exercises for Section 12.6 . . . 950
12.7 Datalog Implementation by BDD's . . . 951
12.7.1 Binary Decision Diagrams . . . 951

12.7.2 Transformations on BDD's . . . 953
12.7.3 Representing Relations by BDD's . . . 954
12.7.4 Relational Operations as BDD Operations . . . 954
12.7.5 Using BDD's for Points-to Analysis . . . 957
12.7.6 Exercises for Section 12.7 . . . 958
12.8 Summary of Chapter 12 . . . 958
12.9 References for Chapter 12 . . . 961

A A Complete Front End **965**
A.1 The Source Language . . . 965
A.2 Main . . . 966
A.3 Lexical Analyzer . . . 967
A.4 Symbol Tables and Types . . . 970
A.5 Intermediate Code for Expressions . . . 971
A.6 Jumping Code for Boolean Expressions . . . 974
A.7 Intermediate Code for Statements . . . 978
A.8 Parser . . . 981
A.9 Creating the Front End . . . 986

B Finding Linearly Independent Solutions **989**

Index **993**