Chapter 29. Selling into Enterprise Markets

Think Lean Analytics only applies to consumer-focused businesses? Think again.

Sure, it's easier to experiment on consumers—there are so many of them out there, and they make decisions irrationally, so you can toy with their emotions. There's no doubt that cloud computing and social media have made it easy to launch something and spread the word without significant upfront investment, and consumer startups are media icons, even fodder for Hollywood.[136] Even business-to-business startups, such as SaaS providers, often target small and medium companies.

But a data-informed approach to business is good for any kind of organization. Plenty of great founders went after big business problems, and got rich doing so. As TechCrunch reporter Alex Williams put it, "While the enterprise can be as boring as hell, the whole goddamn thing is paved with gold."[137] Enterprise-focused startups do have to deal with some unique challenges along the way, which changes the metrics they watch and how they collect them, but it's worth it.

Why Are Enterprise Customers Different?

Let's start with the good news: it's easier to find enterprises to talk to. They're in the phone book. They might have time for coffee. They have budgets. And for many of the people in these organizations, it's part of their *job* to evaluate new solutions, meet with vendors, and share their needs to see if someone can solve them more explicitly. Armed with a decent caffeine allowance, you can talk to actual prospects fairly quickly.

That said, there are plenty of important ways that enterprise sales are different and more difficult than selling to a large, unwashed audience. Venture capitalist

Ben Horowitz was one of the first to burst this bubble:

> *Every day I hear from entrepreneurs, angel investors, and venture capitalists about an exciting new movement called "the consumerization of the enterprise." They tell me how the old, expensive, Rolex-wearing sales forces are a thing of the past and, in the future, companies will "consume" enterprise products proactively like consumers pick up Twitter.*
>
> *But when I talk to the most successful new enterprise companies like WorkDay, Apptio, Jive, Zuora, and Cloudera, they all employ serious and large enterprise sales efforts that usually include expensive people, some of whom indeed wear Rolex watches.*[138]

Big Ticket, High Touch

The one thing that makes enterprise-focused startups different is this: B2C customer development is polling, B2B customer development is a census.

In most cases, enterprise sales involve bigger-ticket items, sold to fewer customers. That means more money from fewer sources. If you're selling a big-ticket item, this changes the game dramatically. For starters, you can afford to talk to every customer. The high sale price offsets the cost of a direct sales approach, particularly in the early stages of the sale.

The small number of initial users makes an even bigger difference. You aren't talking to a sample of 30 people as a proxy for the market at large. Instead, you're talking to 30 companies who may well become your first 30 customers.

Much of analytics is about trying to understand large amounts of information so you can get a better grasp of underlying patterns and act on them. But in the early stages of a B2B startup, there aren't patterns—there are just customers.

- You can pick up the phone and call them right away.
- They'll call you and tell you what they want.
- You can get in a room with them.
- You can't test something on a statistically significant sample of the population and write it off if the test fails—you'll lose customers.

Formality

Enterprise buyers tend to be more regulated. They can't make decisions on gut or emotion—or rather, they can, but it has to be justified with a business case.

Big companies are often public companies with checks and balances. The person who pays for the product (finance) isn't the person who uses it (the line of business). Understanding this dichotomy is critical for product development and sales. Initially, you may target early adopters, where the buyer is much closer to the user (they may be the same person at this point), but as you move past early adopters, the buyer and user diverge.

Companies have formal structure for good reasons. It helps prevent corruption, and makes auditing possible. But that structure gets in the way of understanding things. Your contact at a company may be a proponent, but someone else in the organization may be a detractor, or have a concern of which you're not aware. This is one of the reasons direct sales is common in early stages: it lets you navigate the bureaucracy and understand the part of the sales process that's hidden to outsiders.

Legacy Products

Consumers can ditch their old product on a whim. Small businesses can migrate fairly easily, as the recent exodus to cloud-based software demonstrates. Large companies, on the other hand, have a significant capital investment in the past which must be properly depreciated. They also have a significant political investment in past decisions, and often this is the strongest opposition to change.

Most organizations of any real size have developed their own software and processes, and they expect you to adapt to them. They won't change how they work: change is hard, and retraining is a cost. This can increase your deployment costs, because you have to integrate with what's already in place. It also means your products must be more configurable and adaptable, which translates into more complexity and less ease of use.

Incumbents

Those legacy issues are part of another problem—incumbents. If you're trying to disrupt or replace something, you'll have to convince the organization that you're better, despite the efforts of an existing solution. Organizations are averse to change, and love the status quo. If you're trying to sell to them, and your product is still in the early stages of the technology adoption cycle, you're

penalized simply for being new. *Consumers love novelty; businesses just call it risk.*

This also means incumbent vendors can stall your sale significantly if they get wind of what you're planning to do just by claiming that they're going to do it too. They can step on your oxygen hose by promising something—then rescind the promise once you're dead.

Of course, big, slow incumbents have plenty of weaknesses. New entrants can disrupt their market simply by being easier to adopt, because they require no training. A decade ago, the only people who knew what a "feed" was were stock traders connecting to Bloomberg terminals; today, everyone who's used Facebook or Twitter is familiar with feeds. They don't need to be trained.

Simplicity isn't just an attribute of enterprise disruption—it's the price of entry. DJ Patil, data scientist in residence at Greylock and former head of product at LinkedIn, calls this the Zero Overhead Principle:

> *A central theme to this new wave of innovation is the application of core product tenets from the consumer space to the enterprise. In particular, a universal lesson that I keep sharing with all entrepreneurs building for the enterprise is the Zero Overhead Principle: no feature may add training costs to the user.*[139]

Slower Cycle Time

Lean Startup models work because they empower you to learn quickly and iteratively. It's hard to achieve speed when your customer moves sluggishly and carefully, so the slower cycle time of your target market makes it tough to iterate quickly. This is a key reason why many of the early Lean Startup success stories have come from consumer-focused businesses.

The rise of the SaaS market changes this, because it's relatively easy to alter functionality without the market's permission. But if you're selling traditional enterprise software, or delivery trucks, or shredders, you're not going to learn and iterate as quickly as you would from consumers. Of course, your competitors aren't either. You don't need to be fast—just faster than everyone else.

Rationality (and Lack of Imagination)

Not all companies fit the stereotype of the big, slow, late-adopter customer, but

risk aversion is real. Because enterprise buyers can't take the risks consumers can, they limit their own thinking. They demand proof that something will work before they try it out, which means great ideas can often become mired in business cases, return-on-investment analyses, and total-cost-of-ownership spreadsheets.

This rationality is warranted. In 2005, IEEE (Institute of Electrical and Electronics Engineers) committee chair Robert N. Charette estimated that of the $1 trillion spent on software each year, 5–15% would be abandoned before or shortly after delivery, and much of the rest would be late or suffer huge budget overruns.[140] A similar study by PM Solutions estimates that 37% of IT projects are at risk.[141]

Because companies are full of people—for many of whom their job is just a job—their priority is to minimize the chance of them making a mistake even if the organization as a whole might suffer in the long term. It's hard to inspire an organization if its employees are busy wondering whether the changes you promise will cost them their jobs.

This is an unnecessarily bleak view of the world.

For all these reasons, most B2B-focused startups consist of two people: a domain expert and a disruption expert.

- The domain expert knows the industry and the problem domain. He has a Rolodex and can act as a proxy for customers in the early stages of product definition. Often this person is from the line of business, and has a marketing, sales, or business development role.

- The disruption expert knows the technology that will produce a change on which the startup can capitalize. She can see beyond the current model and understand what an industry will look like after the shift, and brings the novel approach to the existing market. This is usually the technologist.

The Enterprise Startup Lifecycle

Startups begin in many ways. Over the years, however, we've seen a recurring pattern in how B2B startups grow. It usually happens in one of three ways:

The enterprise pivot

In this pattern, the company creates a popular consumer product, then pivots to tackle the enterprise. This is what Dropbox did, and to some extent it's the way BlackBerry circumvented enterprise IT by targeting mercenary salespeople. It's not trivial, though: enterprises have very different expectations and concerns from consumers.

Copy and rebuild

Another approach is to take a consumer idea and make it enterprise-ready. Yammer did this when it rebuilt Facebook's status update model and copied Facebook's feed interface.

Disrupt an existing problem

There are plenty of disruptions that happen to an industry, from the advent of mobile data, to the Internet of Things,[142] to the adoption of the fax machine, to location-aware applications. Any of them can offer a big enough advantage to make it worth discarding the old way of doing things. Taleo did this to the traditional business of human resources management.

Inspiration

Many of the enterprise startups we've talked to began with a basic idea, often hatched within the ecosystem they wanted to disrupt. That's because domain knowledge is essential. Important elements of how a business works—particularly back-office operations—are hidden from the outside world. It's only by being an insider that the bottlenecks become painfully obvious.

Take the founders of Taleo. They left enterprise requirements planning (ERP) heavyweight BAAN to bring talent management tools to the enterprise. They had seen that the big challenges of ERP were integration and deployment, and they'd realized that the Web was how many organizations connect with candidates. They also saw that talent management, both before and after hires were made, was increasingly data-driven.

Many of their realizations came from seeing technology trends. But the founders' fundamental knowledge of the HR industry came from their time at BAAN. Clearly, it worked out well: in February 2012, Oracle acquired Taleo for

$1.9 billion.

That doesn't mean the founding team *must* include an insider—but it helps. Remember, though, insiders still need to "get out of the building" and validate their assumptions; not doing so because of existing domain expertise can be disastrous.

Let's look at how the five stages of the Lean Analytics framework apply to a B2B-focused company. Figure 29-1 shows what a B2B company needs to do at each stage, as well as what risks it should fear.

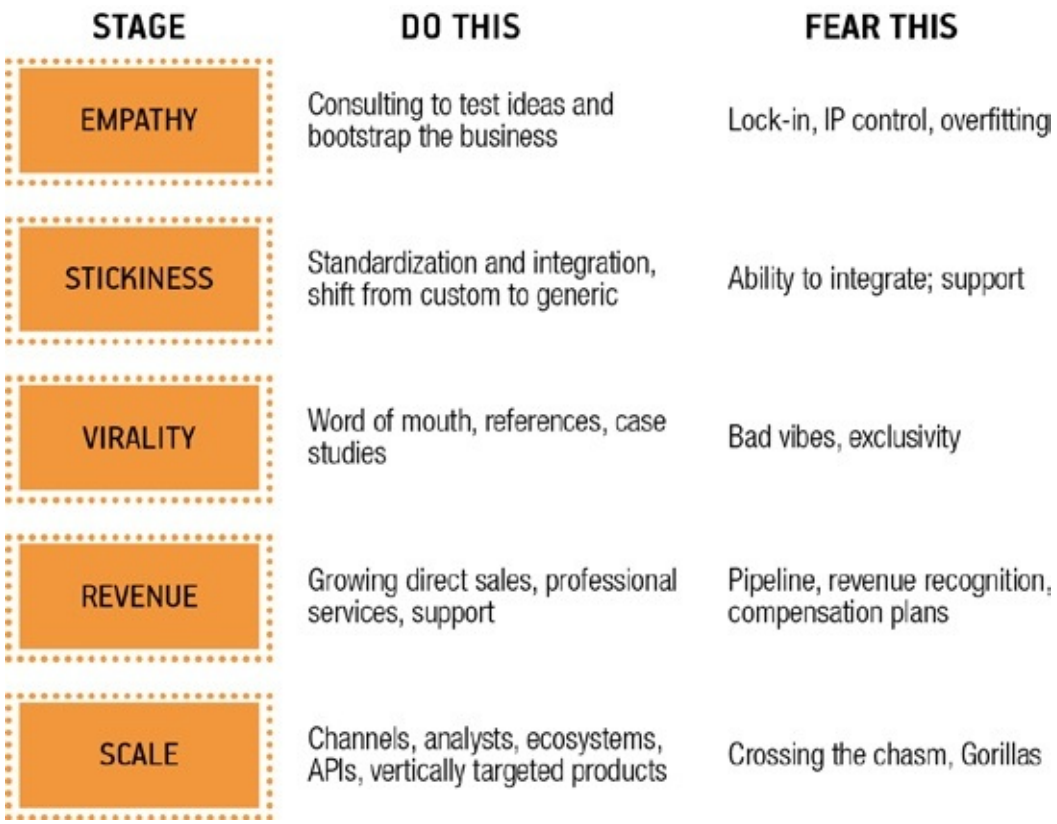

Figure 29-1. The Lean Analytics stages when you're selling to enterprises

Empathy: Consulting and Segmentation

Many bootstrapped startups begin their lives as consulting organizations. Consulting is a good way to discover customer needs, and it helps pay the bills. It also gives you a way to test out your early ideas, because while every customer has needs, the only needs you can build a business on are those that are

consistent across a reasonably large, addressable market.

Having said that, consulting companies struggle a great deal to transition from service providers to product companies because they need to, at some point, abandon service revenues and focus on the product. That transition can be extremely painful—from a cash flow perspective—and most service providers don't make the jump.

It's also necessary to "burn the boats" of the services business to ensure that you commit to the product. After all, you're going to neglect some of your most-loved customers in order to deliver a product the general market wants instead, and it'll be tempting to do custom work to keep them happy. You can't run a product and a services business concurrently. Even IBM had to split itself in two; what makes you think you can do it as a fledgling startup?

How Coradiant Found a Market

Coradiant, a maker of web performance equipment, started in 1997 as Networkshop, and was acquired by BMC Software in April 2011.[143] Initially it was an IT infrastructure consulting firm that wrote studies on performance, availability, and web technologies like SSL.[144] Soon, however, enterprises and startups approached the company seeking help with their deployments. These customers needed several pieces of costly network infrastructure—a pair of load balancers, firewalls, crypto accelerators, switches, routers, and related monitoring tools that, together, cost up to $500,000 and handled 100 megabits per second (Mbps) of traffic. But these companies needed only a fraction of that capacity.

Networkshop built a virtualized front-end infrastructure that customers could buy one Mbps at a time. It deployed this in a single data center in one city, and offered fractional capacity to customers in that data center. The economics were good: once the infrastructure had exceeded 35% utilization by customers, every additional dollar went straight to the bottom line.

Armed with this example, Networkshop changed its name to Coradiant and closed Series A funding, using the proceeds to deploy similar "pods" of infrastructure in data centers throughout North America. Wrapping this in support services, the company joined firms like LoudCloud and SiteSmith in the growing managed service provider (MSP) business.

growing managed service provider (MSP) business.

Within a few years, however, the data center owners with whom Coradiant had colocated realized that they needed to make more money from their facilities. To increase their revenue per square foot, they started offering competing services. The Coradiant founders had a decision to make: either compete head-to-head with the very same data centers in which their customers were hosted—a bad idea—or pivot to a new model that didn't need the data center owners' permission.

Coradiant had built a monitoring service (called OutSight) to help manage customers' infrastructure and measure performance. In the summer of 2003, the company scaled back dramatically, laying off most operational staff and hiring developers and architects who focused on building an appliance version of this technology. The new product, dubbed TrueSight, launched in 2004, and this time, Coradiant didn't need the data center owners' permission to be deployed.

Some of Coradiant's MSP customers became TrueSight users, quickly building a stable of reference-worthy household names. The initial version of TrueSight contained only basic features—most reporting, for example, was done by exporting information into Excel. But Coradiant had an extremely hands-on pre- and post-sales engineering team that worked closely with early customers. Once the company saw what kinds of reports customers made, and how they used the appliance, it incorporated those into later versions.

Coradiant didn't use channel sales until the product was relatively mature. The direct contact helped provide frequent feedback from the field. The company also held user conferences twice a year to hear how people were using the product, which led it into new directions such as real-time visualization and data export for vulnerability detection.

Ultimately, the consulting heritage gave Coradiant insight into the needs of a target market. The initial product offering was based on the sharing of IT infrastructure, amortizing the cost of networking components across many customers. That service, in turn, helped the firm learn what features customers needed from a monitoring product, and ultimately led it to build the product for which it was acquired.

Summary

- Coradiant started selling managed services, but a major market shift changed

the dynamics of the market significantly.

- The company found that its unique value was a subset of the managed services offered that looked at users' experience on a website.
- Customers wanted this functionality as an appliance rather than a service.

Analytics Lessons Learned

Sometimes, environmental changes such as legislation or competition mean that validated business assumptions are no longer true. When that happens, look at what your core value proposition is and see if you can sell it to a different market or in a different way that overcomes those changes—in this case, keeping only a subset of a service and delivering it as an appliance.

Launching a startup as a consultancy has its risks. It's easy to get trapped in consulting. As the business grows, you'll want to keep customers happy, and won't have the cycles to dedicate to building the product or service you want. Many startups have lost sight of their initial plan and are now consulting firms—some of them happily. But they don't meet Paul Graham's test for scalable, repeatable, rapid growth. They're not startups.

What's more, in order to make the shift from consultancy to startup, you first need to test whether your existing customers' demands are applicable to a broader audience. Doing so may violate privacy agreements you have with your customers, so you need to finesse customer development. Your existing clients may feel that a standardized product you plan to offer will be less tailored to their needs; you need to convince them that a standard product is in fact better for them, because the cost of building future versions will be shared among many buyers.

Once you've found the problem you're going to fix, and have verified that the solution will work with your prospects and clients, you need to segment them. Not all clients are identical, so it's smart to pick a geographic region, a particular vertical, or customers who belong to just one of your sales teams. That way, you can give those early adopters better attention and limit the impact of failure.

Imagine, for example, that you're building a hiring management tool. The way that a legal firm finds and retains candidates is very different from the way a fast-food restaurant does it. Trying to build a single tool for them—particularly

at the outset—is a bad idea. Everything from the number of interviews, to the qualifications needed, to the number of years someone stays with the company will be different. Differences mean customization and parameters, which increase complexity, and violates DJ Patil's Zero Overhead Principle.

Stickiness: Standardization and Integration

Once you know the need and have identified your initial segments, you have to standardize the product. With some products, it's possible to sell before building. Instead of an MVP, you may have a prototype, or a set of specifications for which the prospect will commit to paying on delivery. This pipeline of conditional purchases reduces the cost of fundraising, because it increases the chances of success.

In the B2C world, startups worry less about "Can I build it?" and more about "Will anyone care?" In the enterprise market, the risk is more, "Will it integrate?" Integration with existing tools, processes, and environments is the most likely source of problems, and you'll wind up customizing for clients—which undermines the standardization you fought so hard to achieve earlier.

Managing this tension between customization and standardization is one of the biggest challenges of an early-stage enterprise startup. If you can't get the client's users to try the product, you're doomed. And while your technology might work, if it doesn't properly integrate with legacy systems, it'll be seen as your fault, not theirs.

Virality: Word of Mouth, Referrals, and References

Assuming you've successfully sold the standardized product to an initial market segment, you'll need to grow. Because enterprises don't trust newcomers, you'll rely heavily on referrals and word-of-mouth marketing. You'll make case studies from early successes, and ask satisfied users to handle phone calls from new prospects.

Referrals and references are critical to this stage of growth. A couple of household names as customers are priceless. Enterprise-focused vendors will often provide discounts in exchange for case studies.

Revenue: Direct Sales and Support

Revenue: Direct Sales and Support

With the pipeline growing and revenue coming in, you'll worry about cash flow and commission structures for your direct sales team. To know if you have a sustainable business, you'll also look at support costs, churn, trouble tickets, and other indicators of ongoing business costs to learn just how much a particular customer contributes to the bottom line. If the operating margin is bad, it will have a significant drag on profitability.

Feedback from the sales team and the support group is critical at this point, because it indicates whether your initial success is genuine, or simply a case of prospects buying into the story you're telling (which won't be sustainable in the longer term). Zach Nies, Chief Technologist at Rally Software says, "This is absolutely critical for startups, because they have a huge advantage here. In most incumbents, the product development team is so far removed from the field and customers that they have no sense of trends in the market. Often startups will know a lot more about an incumbent's customers than the incumbent does."

Scale: Channel Sales, Efficiencies, and Ecosystems

In the final stages of an enterprise-focused startup, you'll emphasize scaling. You may have channel sales through value-added resellers and distributors. You'll also have an ecosystem of analysts, developers, APIs (application programming interfaces) and platforms, partners, and competitors that will define and refine the market. These are all good indicators that companies will keep using you, because they're investing in processes, vendor relationships, and technology that will make it harder for them to leave you. Scaling an enterprise software company takes years to accomplish. Zach estimates that it can be as long as 5 to 10 years before a company selling into the enterprise has established and validated channels, and mastered its sales processes.

So What Metrics Matter?

Just as there are plenty of parallels between the way B2C and B2B startups grow, so many of the metrics we've seen for consumer-focused companies apply equally well to enterprise-focused ones. But there are a few metrics that you'll want to consider that apply more to enterprise startups.

Ease of Customer Engagement and Feedback

As you're talking to customers, how easy is it to get meetings with them? If you plan to use a direct sales organization later on, this is an early indicator of what it'll be like to sell the product.

Pipeline for Initial Releases, Betas, and Proof-of-Concept Trials

As you start to sign up prospects, you'll track the usual sales metrics. Unlike B2C platforms where you're looking at subscription and engagement, if you're selling a big-ticket, long-term item, you're looking at contracts. While you may not have recognizable revenue, you'll have lead volume and bookings to analyze, and these should give you an understanding of the cost of sales once the product has launched.

It's important—right from the very beginning—that you articulate the stages of your sales funnel and the conversion rates at each point along the way. The sales cycle needs to be well documented, measured, and understood after the first few sales, to see if you can build a repeatable approach. At that point, you can bring in additional salespeople to increase volume.

Stickiness and Usability

As we've seen, the usability of a disruptive solution is "table stakes" for a new entrant in today's market. Companies expect ease of use, because they didn't have to get trained on Google or Facebook, and thus shouldn't have to get training from you, either. DJ Patil suggests using data to find where the friction is hiding in your usage and adoption. "If you can't measure it, you can't fix it," he says. "Instrument the product to monitor user flows and be able to test new ideas in how to iteratively improve your product."

Integration Costs

In the heat of the moment, it's hard to take notes, but integration plays such a big role in enterprise sales that you have to be disciplined about measuring it. What's the true cost of pre-and post-sales support? How much customization is required? How much training, explaining, and troubleshooting are you doing in order to successfully deliver a product to a customer?

order to successfully deliver a product to a customer?

You need to capture this data early on, because later it's an indicator of whether you've built a startup or just created a highly standardized consulting practice. If you prematurely accelerate the latter, thinking it's the former, supporting an expanded market and a sales channel will crush you. This data can also be used against incumbents in a total-cost-of-ownership analysis.

User Engagement

No matter what you're building, the most important metric is whether people are using it. In an enterprise, however, the buyer is less likely to be the user. That means your contact may be an IT project manager, someone in purchasing, or an executive, but your actual users may be rank-and-file employees with whom you have no contact.

You may also have to refrain from talking to users: it's easy to pop up a survey on a consumer website, but employers may frown upon you using up their employees' precious time to answer your questions.

Simply measuring metrics like "time since last use" will be misleading, too, because users are *paid* to use your tool. They may log in every day because it's their job to do so; that doesn't mean they enjoy it. The real questions are whether they *like* logging in, and whether it makes them more productive. Users have a task they want to accomplish, and your product will thrive if it is the perfect tool for that task. Some marketers advocate analyzing customer needs by the job the customer is trying to get done (known as the "jobs-be-done" approach) rather than by segments.[145]

Get baselines from your clients that apply to their real-world businesses before you deploy. How many orders do they enter a day? How long does it take an employee to get payroll information? How many truck deliveries a day can their warehouse handle? What is the usual call hold time? Once you've deployed, use this information to measure progress, helping your advocates to prove the ROI—and turning it into case studies you can share with other customers.

Disentanglement

As you transition from a high-touch consulting business to a standardized one with less customer interaction, you need to focus on disentanglement. Your goal

is to not have “anchor” customers that represent a disproportionate amount of your revenue or your support calls, because you need to scale.

Put your high-touch customers that you acquired early on into a segment and compare them to the rest of your customers. How do they differ? Do they consume a fair proportion of your support resources? Do their feature requests match those of all your customers and prospects? Don’t ignore the companies that made you who you are—but do realize they’re not in a monogamous relationship with you anymore.

Zach Nies suggests going even further, segmenting customers into three groups. “‘A customers’ are your really big customers who negotiated a big discount and expect the world from you. ‘B customers’ are customers who are fairly low maintenance, didn’t get a big discount, see themselves as partners with you, and provide useful insights. ‘C customers’ cause trouble, are a pain to deal with, and demand things from you that you feel will damage your business,” he explains. “Don’t spend too much time on the A’s—they sound good but aren’t the best for your business. Bring as many Bs on as customers as possible. And try to get your ‘C customers’ to be customers of your competitors.”

Support Costs

Zach’s advice is based on some fundamental truths. In many B2B-focused companies, the top 20% of customers generate 150–300% of profits, while the middle 70% of customers break even, and the lowest 10% of customers reduce 50–200% of profits.[146]

You’ll track support metrics like top-requested features, number of outstanding trouble tickets, post-sales support, call center hold time, and so on. This will indicate where you’re losing money, and whether the product is standardized and stable enough to move into growth and scaling.

Segment this data, too. Figure out who’s costing the most money. Then consider firing them.[147] Once, it was hard to break out individual customer costs, but electronic systems make it possible to assign activities—such as support calls, emails, additional storage, or a truck roll—to individual customers.

You don’t actually have to fire customers, of course. You can simply change their pricing enough to make them profitable or encourage them to leave. This is part of getting your pricing right before you grow the business to a point where

part of getting your pricing right before you grow the business to a point where unprofitable clients can do real damage at scale.

User Groups and Feedback

If your business involves big-ticket sales, you may have few enough customers that you can get many of them in the same room. Informal interaction with existing customers can be a boon to enterprise-focused startups, and resembles the problem and solution validation stages of the Lean Startup process—only rather than validating a solution, you're validating a roadmap. Even with a large number of customers, Zach says, "Identify the real advocates and bring them in for a big hug." He also suggests helping advocates network among themselves, which Rally does on its website.[148]

Successful user-group meetings require considerable preparation. Users will be eager to please—or quick to complain—so results will be polarized. They'll also agree to every feature you suggest. Force them to choose; they can't have everything, and you need to present them with hard alternatives (also known as *discrete choices*).

A lot of work has gone into understanding how people make choices. "A 'discrete' choice," says Berkeley professor Dan McFadden, "is a 'yes/no' decision, or a selection of one alternative from a set of possibilities." His application of discrete choice modeling to estimate the adoption of San Francisco's Bay Area Rapid Transit system—which was under construction at the time of his research—earned him the 2000 Nobel Prize in Economics.[149] One important conclusion from this work is that people find it easier to discard something they don't want than to choose something they do (which feels like commitment), so a series of questions in which they are asked to discard one of two options works well.

The math of choice modeling is complex. There are entire conferences devoted to the subject, and it's widely used in new product development for everything from laundry detergent to cars. But some of the methodologies are instructive. For example, you can get better answers by repeatedly asking your customers to compare two possible feature enhancements and choosing the one they can do without, rather than by simply asking them to rate the possible features on a scale of 1 to 10. You'll do even better if you mix up several attributes in each comparison, regardless of whether a particular combination of attributes makes

sense.

Imagine you're trying to find a new diet food to introduce. You know the attributes that might affect buyers include taste, calories, gluten content, and sustainable ingredients. Simply asking prospects whether taste is more important to them than caloric content is informative. But asking them to make a choice between two discrete offerings—even if those offerings are theoretically impossible—is even better. Would you prefer:

- A delicious, gluten-free, high-calorie candy made with artificial ingredients;
- Or a bland, high-gluten, low-calorie candy of organic origin?

Asking customers to trade off variations of combinations, over and over, dramatically improves prediction accuracy. In fact, this is equivalent to the multivariate testing we've discussed before, applied to surveys and interviews.

As you're designing user events, know what you're hoping to learn and invest in the conversations and experimental design needed to get real answers that you can turn into the right product roadmap.

Pitch Success

You've measured your effectiveness at setting up meetings in the early phases of your startup. It matters later on, when you're about to bring on channels. Your channel partners aren't as clever as you, and you'll need to arm them with collateral and messaging that they can use to close deals without your assistance. If they try to push your product or service and encounter resistance, they'll sell something else. With channels, you seldom get a second chance to make a first impression.

Create marketing tools for your channel and then test them yourselves. Make cold calls with their scripts. Pitch them to new customers. Send out email form letters and test response rates.

This does two things: first, it shows you which script, pitch, or form letter to use (because, after all, everything's an experiment, right?, and second, it gives you a baseline against which to compare channel effectiveness. If a channel partner isn't meeting your baseline, something else is wrong, and you can work to fix it before that partner sours on your product.

If you make channel collateral, tag each piece of collateral with something that identifies the channel. You might use shortened URLs that include a code identifying the partner in PDFs you create, which would let you see which partners' efforts are driving traffic to your site.

Barriers to Exit

As you bring customers on at scale, you want to make them stick around. A vibrant developer ecosystem and a healthy API allow customers to integrate themselves with you, making *you* the incumbent vendor and helping you to counter threats from competitors and new entrants.

Simon Wardley, who studies organizational warfare and evolution for the Leading Edge Forum, points out that companies must prioritize the long list of features customers need. Build too many, and they won't all be profitable; build too few, and you leave the door open to competitors. APIs, he says, offer a solution.[150]

> *All innovations... are a gamble and whilst we can reduce costs we can never eliminate it. The future value of something is inversely proportional to the certainty we have over it; we cannot avoid this information barrier any more than we can reliably predict the future. However, there is a means to maximize our advantage.*
>
> *By making these utility services accessible through APIs, we not only benefit ourselves but we can open up these components to a wider ecosystem. If we can encourage innovation in that wider ecosystem then we do not incur the cost of gambling [and] failure for those new activities. Unfortunately, we do not enjoy the rewards of their success either.*
>
> *Fortunately, the ecosystem provides an early warning mechanism of success (i.e., adoption)...by creating a large enough ecosystem, we can not only encourage a rapid rate of innovation but also leverage that ecosystem to identify success and then either copy (a weak ecosystem approach) or acquire (a strong ecosystem approach) that activity. This is how we maximize our advantage.*

If you have an API, track its usage by clients. Those clients who have a lot of API activity are investing more in extending their relationship with you; those who are inactive could switch vendors more easily. If you have a developer program, examine searches and feature requests to discover what tools your customers want, then find developers to build features you aren't going to create yourself.

The Bottom Line: Startups Are Startups

While enterprise-focused startups must contend with some significant

differences, the fundamental Lean Startup model remains: determine the riskiest part of the business, and find a way of quantifying and mitigating that risk quickly by creating something, measuring the result, and learning from it.

---

[136] In February 2012, The Next Web's Allen Gannett listed the rise of the cloud, the consumerization of technology, and the broad adoption of SaaS delivery models as three catalysts for the rapid expansion of acquisitions in enterprise software.

[137] Williams's reaction after attending a demo day by Acceleprise, an accelerator focused on startups that target enterprise customers; see *http://techcrunch.com/2012/11/09/notes-from-a-startup-night-the-enterprise-can-be-as-boring-as-hell-but-the-whole-goddamn-thing-is-paved-with-gold/*.

[138] *http://bhorowitz.com/2010/11/15/meet-the-new-enterprise-customer-he%E2%80%99s-a-lot-like-the-old-enterprise-customer/*

[139] *http://techcrunch.com/2012/10/05/building-for-the-enterprise-the-zero-overhead-principle-2/*

[140] *http://spectrum.ieee.org/computing/software/why-software-fails/0*

[141] *http://www.zdnet.com/blog/projectfailures/cio-analysis-why-37-percent-of-projects-fail/12565*

[142] *http://en.wikipedia.org/wiki/Internet_of_Things*

[143] Full disclosure: Coradiant was co-founded by Alistair Croll and Eric Packman in 1997 as Networkshop; the name was changed to Coradiant in mid-2000.

[144] *http://www.infosecnews.org/hypermail/9905/1667.html*

[145] *http://hbswk.hbs.edu/item/6496.html*

[146] Robert S. Kaplan and V.G. Naranyanan "Measuring and Managing Customer Profitability," *Journal of Cost Management* (2001), 15, 5–15, cited in Shin, Sudhir, and Yoon, "When to 'Fire' Customers."

[147] Jiwoong Shin, K. Sudhir, and Dae-Hee Yoon, "When to 'Fire' Customers: Customer Cost-Based Pricing," *Management Science*, December 2012 (*http://faculty.som.yale.edu/ksudhir/papers/Customer%20Cost%20Based%20Pricing.pdf*).

[148] *http://www.rallydev.com/community*

[149] *http://elsa.berkeley.edu/~mcfadden/charterday01/charterday_final.pdf*

[150] *http://blog.gardeviance.org/2011/03/ecosystem-wars.html*